Explaining the Development of International Norms: The Humanitarian Turn at the United Nations Security Council

By Richard Hanania

Columbia University

Saltzman Institute of War and Peace Studies

June 28, 2019

Abstract

The UN Security Council (UNSC) has transformed from a body almost exclusively focused on conflict to one that addresses a variety of issues. Unfortunately, we still lack clear understanding of why and when international institutions change their missions. The author argues that while international politics is usually characterized by inertia, shocks to the system, or focal point events, can compel actors to adopt new logics of appropriateness. Since 1945, the end of the Cold War and the signing of the Helsinki Accords stand out as such events. Through unstructured topic modeling, UNSC resolutions divide into the subjects of *War*, *Punitive*, and *Humanitarian*. The topic *Humanitarian* exploded in frequency after the Cold War, and more refined models show that words related to human rights and elections similarly increased after Helsinki. These changes are rapid and occur in the immediate aftermath of focal point events, showing their importance for norm diffusion.

9,968 Words

As he was negotiating the creation of the United Nations Security Council (UNSC), President Roosevelt was of two minds. Despite telling British Foreign Minister Anthony Eden in 1943 that he saw the institution as a means to disarm less powerful countries, FDR occasionally talked of a new world order in which the more developed nations would preside over a relatively egalitarian arrangement and take on the responsibility of protecting human rights. Stalin and Churchill, Roosevelt's negotiating partners, were less ambivalent about what they wanted out of the postwar system. They sought to secure the territories and borders of their respective empires, and believed that the job of the UNSC should simply be to keep the peace between states (Bosco 2009:14–23).

For its first several decades, the UNSC acted in accordance with the vision of Churchill and Stalin. Its resolutions addressed issues that one would have naturally expected to be of concern to the great powers (Moyn 2010:8). Over time, however, something changed. Resolution 2041 (2012), for example, was on its surface about the ongoing occupation of Afghanistan and extended the UN mission in that country. Yet the document also stressed the importance of and put forth concrete steps towards achieving goals such as including women in state decision-making processes and reforming the prison sector. Although the threat of terrorism had originally brought the international community into Afghanistan, the resolutions addressing the ongoing conflict have given the forces in that country a role more in line with the more idealistic vision of Roosevelt.

How and when did this change happen? This article, while accepting the role of historical contingency, presents a general theory of institutional development. The UNSC is chosen as the focus of this analysis because of its central role in the creation of

international law, as it is the only institution able to directly do so on matters of war and peace. Because of this, it possesses unique legitimacy as a focal point of international politics, which gives it an unparalleled ability to effectively engage in norm diffusion (Hurd 2008). Borrowing from game theory, this article argues that major changes in international practice may be compelled by focal point events, or exogenous shocks that can fundamentally change the focus and purposes of an international organization. Using a machine learning algorithm called Latent Dirichlet Allocation, it finds that the UNSC shifted towards more humanitarian topics in the aftermath of two focal point events: the signing of the Helsinki Accords and the collapse of the Soviet Union. Gaining insights into the circumstances surrounding major changes in the focus of the UNSC, along with their timing, is important for understanding the development of international society over the last three-quarters of a century, and can help shed light on processes of norm diffusion and the creation of international law more generally.

Changing Values in International Society

Few questions in international relations are as fundamental as those that ask about the extent to which ideas about appropriate behavior influence conduct and, in contrast, how much relations between states are driven by narrowly defined interests. Those who believe in the power of normative constraints have in recent years produced a series of powerful works that reveal how the values of international society have changed over time (Pinker 2011, 2019; Hathaway and Shapiro 2017; Goldstein 2012; Gat 2006; Mueller 2009). Industrial and post-industrial societies see few war deaths per capita than our hunter-gatherer ancestors, and, on a more compressed timescale, the second half of the twentieth century and beyond has been particularly peaceful, with the rate and

deadliness of wars reaching all-time lows over the last few decades (Pinker 2011:295–307).

Constructivists have sought to explain at how change happens, and they begin by noting that each international actor has an identity that shapes its actions. An identity includes information such as proper goals, acceptable means to achieve those goals, and what logic of appropriateness to follow (March and Olsen 2011; Checkel 2005; Müller 2004). Collectively held ideas about appropriate behavior consistent with an identity are referred to as norms (Finnemore 1993). For decades now, the idea that norms matter in international relations has been accepted by most scholars in the field, and this view is supported by historical research into changing practices over time (Checkel 1997).

This field of inquiry has also led scholars to investigate the related question of where norms in the international system come from. While this has proved more difficult to answer than the question of whether norms are influential, certain ideas have gained currency. Historical study reveals the prominent role played by norm entrepreneurs, who are unsurprisingly most likely to succeed when they have the backing of a powerful state, as was the case with the British ending the international slave trade and the United States establishing the post-1945 order (Ikenberry 2009; Kaufmann and Pape 1999). Powerful actors are by definition able to use military and economic coercion to achieve their preferred outcomes. Eventually, actors are socialized to the point that they internalize a norm, and it continues to influence behavior even without pressure from the moral entrepreneur (Finnemore and Sikkink 1998).

In addition to processes based on power disparities, scholars also emphasize the importance of persuasion and social influence in international politics (Johnstone 2011;

Johnston 2001). Persuasion is fundamental to political communication, and is most likely to be effective when its object can be shown to be behaving in a way that is incompatible with an already accepted identity (Legro 1997). Social influence applies when an actor such as a state is respected for its perceived accomplishment, skill, or virtue, and because of these qualities is in a position to nudge others into accepting its view of international relations (Gheciu 2005). The two kinds of influence are self-reinforcing and no strict line can be drawn between them; for example, those with social prestige tend to be more persuasive, while being persuasive can bring prestige.

In addition to occuring between states, socialization also takes place in the context of international organizations. Partly, this is because such institutions act as forums in which individuals can influence one another. This is a subtle and less conscious process, through which the act of being imbedded within an organization creates a new identity related to membership in that institution (Lewis 2005; Checkel 2005). Representatives of states, along with the permanent bureaucracy of organizations, follow a script circumscribing a range of action appropriate for them based on this identity, in addition to state-centric and other forms of identity.

There are two main ways that international organizations can spread norms to the rest of the international community. First, in areas where they exercise power, they can take *direct action*, or make policy that incorporates certain norms. International organizations often have their own budgets and mandates from states, and control over resources and their distribution can directly shape the behavior of others (Hurd 2008). Once a state accepts a new practice in order to comply with a regulation from abroad, that norm can become embedded in domestic politics through the activities of bureaucracies

and institutions created in order to ensure compliance (Cortell and Davis 2000; Park 2005). International organizations can also directly sanction behavior that violates widely accepted norms through refusing to provide resources, placing certain conditions on their distribution, or even by simply deciding who gets to take advantage of mechanisms designed to facilitate mutually beneficial cooperation (Hathaway and Shapiro 2011).

In addition to taking direct action, international institutions can spread norms through statements and declarations that are aspirational in nature, or *normative declaration*. This can contribute to the creation of new logics of appropriateness that can change the behavior of domestic and international actors (Finnemore 1993). For this reason, NGOs and governments fight to control the agenda and influence the pronouncements of international institutions, expending tangible resources in order to do so, even when there is no direct payoff or lawmaking power at stake such as in the case of the UN General Assembly (Hurd 2008:112–18).

The Security Council as a Diffusor of Norms

The UNSC is unique among international institutions in its prestige and ability to facilitate action. Indeed, "[t]he ability [of the Council] to mobilize massive coercive resources is unprecedented among international organizations, and almost all states in the system have consented to it in a highly public way." (Hurd 2008:30) The UNSC spreads norms through both methods described above. With regards to direct action, recent decades have seen a rise in UN-created peacekeeping initiatives, and they appear to reduce the likelihood of developing nations reverting to civil war (Fortna 2008). Therefore, by undertaking such missions, the affirmative answer to the question of whether the Security Council can make international law meets even the most stringent

requirements: the institution sends armed personnel into certain countries and forces individuals to comply with the contents of its pronouncements.[1] In authorizing missions in countries such as Afghanistan and Sierra Leone, the Security Council promotes values such as good government, including women in decision-making processes, and a free and open press. Similarly, it sanctions the outcasts of international society such as terrorists and those who produce prohibited weapons, and such punishments tend to be enforced by states and other international institutions (Hathaway and Shapiro 2011; Reinisch 2001).

The Security Council also spreads norms through normative declaration, a practice that has increased over the last few decades. While the effect of direct action can be traced quite easily, how can we likewise be confident that normative declarations have had their intended effect? There are two lines of evidence that suggest that even when it is simply making pronouncements that it does not intend to directly enforce, the Security Council can be an influential force in international politics. First, as mentioned above, states and other actors such as NGOs invest tangible resources in trying to influence the Security Council, including lobbying for seats (Hurd 2008:111–36). This even extends to controlling the agenda, or whether certain issues will or will not be discussed, regardless of whether those issues have a chance of finding their way into resolutions. NGOs likewise push for resolutions that on their surface appear to have no direct effects in terms of the expenditure of resources (Shepherd 2008:386–92; True-Frost 2007:156–57). A particularly vivid recent example of how even symbolic UN votes matter is the late

[1] This view has been referred to as the "Brute Force Objection" to international law (Hathaway and Shapiro 2011:267–68).

2017 controversy of whether to declare Jerusalem the capital of Israel, which for months occupied the leaders of the major powers (Liebermann 2017). Clearly, relevant actors behave as if the pronouncements of the UN matter, and, given their levels of investment in the issues, it would be difficult to argue that they are mistaken in that belief.

The UNSC has power to not only allow the use of force on the part of states, but to require it, making it the global institution with perhaps the greatest ability to diffuse norms. While dealing with international conflicts was its main goal at creation, the UNSC has become more concerned with issues that in previous decades had seen as purely domestic in nature (Joachim 2003). We can call this change in the focus of the UNSC the "Humanitarian Turn," which can be defined as the increasing tendency to address issues that pertain to the rights of individuals and domestic arrangements within countries. One may conceptualize this as a shift away from a focus on protecting states to a focus on protecting individuals (Teitel 2001).

This means that not only has the nature of peacekeeping changed, but we have also seen the passing of what have been called thematic issues peace and security resolutions ("TIPS Resolutions"), which "address issues as diverse as women's equality, HIV/AIDS, and children's rights, and urge all—not just select—states to take action." (True-Frost 2007:116) In other words, these documents are not tied to any particular conflict, nor even the original mission of the UN. They also fail to include any enforcement mechanism, in contrast to humanitarian principles written into peacekeeping resolutions. Nonetheless, a series of case studies show how UN resolutions passed through the Security Council and General Assembly have influenced domestic actors

seeking political change in a variety of countries across the world (Risse and Sikkink 1999; Cohn 2008).

Why did the Humanitarian Turn happen when it did? Theories tend to come in two forms. First, some ground the Humanitarian Turn to specific events within the Security Council itself. For example, in 1991, Secretary General Pérez de Cuéllar issued a report calling on the United Nations to take a stand against human rights violations, and scholars have marked this as an important milestone (Mertus 2009:100–01). Sometimes explanations also account for the role of outside actors. In recent decades the UNSC has seemed to be particularly sensitive to how war affects women and girls, a change that is often traced to Resolution 1325 (2000) (Tryggestad 2009). Yet while norm entrepreneurs in the form of NGOs have helped bring attention to the issue, such groups have been seeking to influence the international community for over a century and a half. Why then, have they only had success in the last few decades (Park 2005:116)? If this is to be attributed to socialization undertaken by NGOs, one has to ask whether they have suddenly become more persuasive in recent years (O'Brein 2000). A much more likely scenario is that there has been a psychological change among those who are being socialized, or the members of the Security Council. Furthermore, the rise of TIPS resolutions and an increase in concern with human rights in peacekeeping missions seem to call for a more holistic explanation. The Security Council became more interested in the rights of women and girls around the same time it was deciding to start taking on environmental issues and protecting other vulnerable groups. As it is highly unlikely that NGOs working on these various issues all became more effective in their work at around the same time, we need a more general theory to account for the Humanitarian Turn.

Indeed, scholars have not failed to notice that the rise in humanitarian pronouncements at the Security Council came after the end of the Cold War. Thus, according to one theory, the great superpower conflict of the second half of the twentieth century prevented the Council from cooperating as much as it could have (Mertus 2010:8–9; Schweigman 2001:287–88). The problem with this view is that much of the rise in Council activity seems to be related to topics that are among the least relevant to Cold War rivalries–e.g., women, children, the protection of civilians, and human trafficking. As will be shown below, the portion of Council work devoted to conflict itself, in contrast, has remained relatively constant over time. During the Cold War, the superpowers were able to agree on broad principles of human rights, as reflected in the Helsinki Accords (1975) (Thomas 2001:81–88). Yet there was no attempt to work such principles into the daily activities of the Security Council itself, even in cases where the superpowers could agree on a peacekeeping mission more generally. For example, the UN Peacekeeping Force in Cyprus was formed in 1964, and although the Security Council passed dozens of resolutions addressing that mission and the conflict itself, in those texts we see very little discussion of the humanitarian concerns that would later become the hallmarks of resolutions related to ongoing conflicts. If the permanent members of the Council were able to achieve consensus on the creation of this and other missions in the first place, why could they not agree to include in their resolutions language pertaining to less controversial matters of human rights? It seems highly likely that the end of the Cold War is in some ways related to the changes at the Security Council, but the exact mechanisms involved remain to be explored.

Overcoming Inertia

This section argues that major changes in international politics are can be created through highly visible exogenous shocks to the system, or focal point events. O'Neill introduces the concept of focal point outcomes in international politics, which are game outcomes that individuals expect to arrive at due to their beliefs about what others believe about their own "consideration of extra-game-theoretic factors." (O'Neill 2005:45–50) In other words, second- and higher-order beliefs about other actors make certain outcomes more likely than others. Similarly, in international politics there are events that are important because everyone believes that everyone else believes that they will be seen as important. Actors are therefore more likely to change their behavior under such circumstances, as the recent shock causes them to rethink "the way things are done" and come to new understandings about the role of institutions (Widmaier, Blyth, and Seabrooke 2007).

Socialization between states is likely to take hold when the socialized country faces a crisis that calls into question the legitimacy of those in power (Ikenberry and Kupchan 1990; Kim and Sharman 2014). International institutions in particular can be quite resistant to change, particularly if their practices become routinized and their attention to rules and procedures prevents them from effectively facing new challenges (Barnet and Finnemore 1999; Shanks, Jacobson, and Kaplan 1996; Hopf 2010). Thus, just as we see exogenous shocks being the catalysts for change in domestic politics, we may expect to see a similar dynamic develop within international institutions.

Focal point events depend on what economists call common knowledge, or knowing that everyone else knows the relevant information (Simler and Hanson

2017:61–62). Common knowledge is of fundamental importance to international law and international relations more generally. For example, the justification for punishing an aggressor lies not only in the first-order effect of not punishing, or letting everyone know that the actor got away it. The second-order effect is also important, as each actor knowing that every other actor is more likely to believe that aggression pays, which can make states more afraid of one another and exacerbate the security dilemma (Hathaway and Shapiro 2017:345–49). Compliance with international law depends thus depends on higher-order beliefs.

Actors and institutions are therefore unlikely to develop a new logic of appropriateness based on information that only they themselves know. The United Nations depends on states for funding and personnel, and states in turn look to it for legitimacy (Hurd 2008). An international organization is less likely to try to spread norms if it has no way of enforcing them, while a state is less likely to carry out enforcement if doing so will be seen as illegitimate. A focal point event can modify each side's expectations about what others will do, thereby changing its own behavior. The same applies to other actors such as NGOs, and how they interact with other agents in the international system. NGOs greatly increased their activity after the collapse of the Berlin Wall, likely in part because donors and others saw more potential for them to be effective (Joachim 2013).

One may very well argue that norms against human rights violations existed well before the end of the Cold War. Why then, did, the Humanitarian Turn take place as late as it did? In order to answer this question, it is useful to divide the socialization process that leads to norm diffusion in the international arena into two components. First, the

socialized actors must be convinced that a specific outcome–e.g., protecting civilians, promoting women's rights, spreading Judeo-Christian values–is normatively desirable. Second, relevant actors need to come to accept that the role that they have taken on in spreading the norm is an appropriate one given their position. The first requirement does not necessarily lead to the second; in fact, it usually does not (Braumoeller 2010). The Humanitarian Turn did not involve member states adopting ideals that were not yet present in international society. Rather, they changed their ideas about what the proper role of the UNSC was. While western governments in particular have worked to spread human rights for centuries (Bass 2008; Kaufmann and Pape 1999), it took focal point events in order to reconceptualize the rights and duties of the Security Council and lead it to accept new roles for itself in the midst of its continuing focus on preventing and resolving conflicts.

Fortunately, we can pin down exactly when the Humanitarian Turn emerged through text analytic methods. Given that new logics of appropriateness can result from focal point shocks, two candidates emerge as possible catalysts for the Humanitarian Turn. The Helsinki Accords (1975) are considered a breakthrough in human rights law and motivated reform movements across the Soviet Union and its satellites (Thomas 2001:206). Before that time, communist power had seemed to be permanently entrenched, as highlighted by a crackdown on political activity in the years immediately preceding it. The Helsinki Accords were a catalyst for the mobilization of transnational networks in Eastern bloc countries such as Poland and Czechoslovakia, and motivated other states to begin putting pressure on the USSR by emphasizing human rights in the course of diplomacy centered around other issues. These agreements signed were

therefore "an essential feature of the crystallization of international human rights consciousness in 1976–77." (Moyn 2010:150)

The signing and ratification of the Helsinki Accords was not an isolated event, but rather part of a movement that throughout the 1970s for the first time made human rights a priority in the normal course of international politics (Moyn 2010:148–55; Denhert 2013). This plausibly helped lead to the end of the Cold War, the second major shock of the second half of the twentieth century. The events of late 1991 meant that, practically overnight, the great ideological conflict of the second half of the twentieth century was over, the largest empire in the world had collapsed, and the international system went from a bipolar structure to a unipolar one. For the reasons discussed above, we should see major changes at the Security Council around the time of the two focal point events discussed above.

Methodology

In recent years, unsupervised topic modeling has emerged as a powerful tool to find evidence for theories that had previously been difficult to validate (Grimmer and Stewart 2013). The last few decades have seen a series of works that have enriched our understanding of legal and political texts, particularly parliamentary speeches, treaties, constitutions, and documents produced by and submitted to courts (Spirling 2012; Evans et al. 2007; Law 2016; Diermier et al. 2012; Slapin and Proksch 2008). Among the most advanced and fruitful ways that unsupervised text analytic models have been used is Latent Dirichlet Allocation (LDA) (Blei, Ng, and Jordan 2003). Analyzing UNSC resolutions is well suited for this method, as it provides constant output in a standardized format that is centered around related but nonetheless discreet topics (Grimmer and

Stewart 2013). LDA assumes that each word in a document *d*, is generated in the following manner. First, choose θ_d, or the topic distribution for that document, from an underlying Dirichlet distribution of topic probabilities. Then, for each word *n*, or w_n one chooses a word from the underlying topic probability distribution of topic z_n. Figure 1 represents this process (Blei, Ng, and Jordan 2003). Plates indicate a loop, meaning that for all documents (*D*), all the words in that document (N_d) are created through the same generative process. The shaded circle represents all that can be directly observed, while each topic-distribution θ and z must be inferred from the data (see Online Appendix).

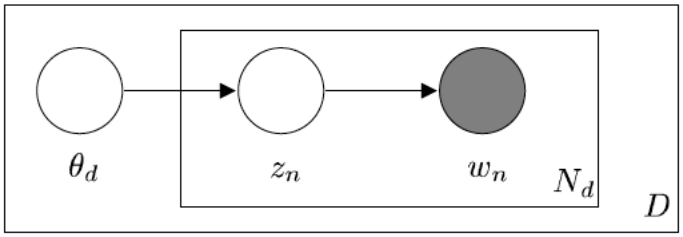

Figure 1: LDA Model

For this paper, the corpus is all 2,397 UNSC Resolutions between the founding of the institution in 1946 up to 2017. As a first step, the names of countries, certain international institutions, and common phrases are each combined into single terms. Among these are terms such as "Democratic People's Republic of Korea," "United Nations," "Western Sahara," "Secretary General," "al-Qaida," and "Sierra Leone." Certain phrases that are common parts of country names are removed, such as "Kingdom of" and "Republic of."[2] Then terms are stemmed and de-capitalized, and stop words, punctuation, and numbers are all removed, before being put into a document-term matrix.

[2] See Online Appendix for more details on this process.

Much discussion revolves around how to choose the number of topics *k*, with some adopting what has been called a Hierarchical Dirichlet Process (HDP) (Teh et al. 2005). However, this process has been criticized for finding topics that group together terms that appear unrelated (Deveaud, SanJuan, and Bellot 2014; Xie and Passonneau 2012). It is also computationally expensive and adds a layer of complexity that can make interpretation difficult. This article therefore uses more recently derived methods designed in order to find the ideal number of topics in an LDA model, relying on Kullback-Leibler divergence, density, Jensen-Shannon divergence, or maximum likelihood estimations (Arun et al. 2010; Cao et al. 2009; Deveaud, SanJuan, and Bellot 2014; Griffiths and Steyvers 2004). Although the results differ in particulars, they generally involve finding a *k* value that achieves either maximum dissimilarity or minimum similarity between topics. In Figure 2, I apply the four methods to the corpus, at various *k* values as indicated by the points on the on the graphs, with the legends indicating the method used.

Figure 2. Model Selection

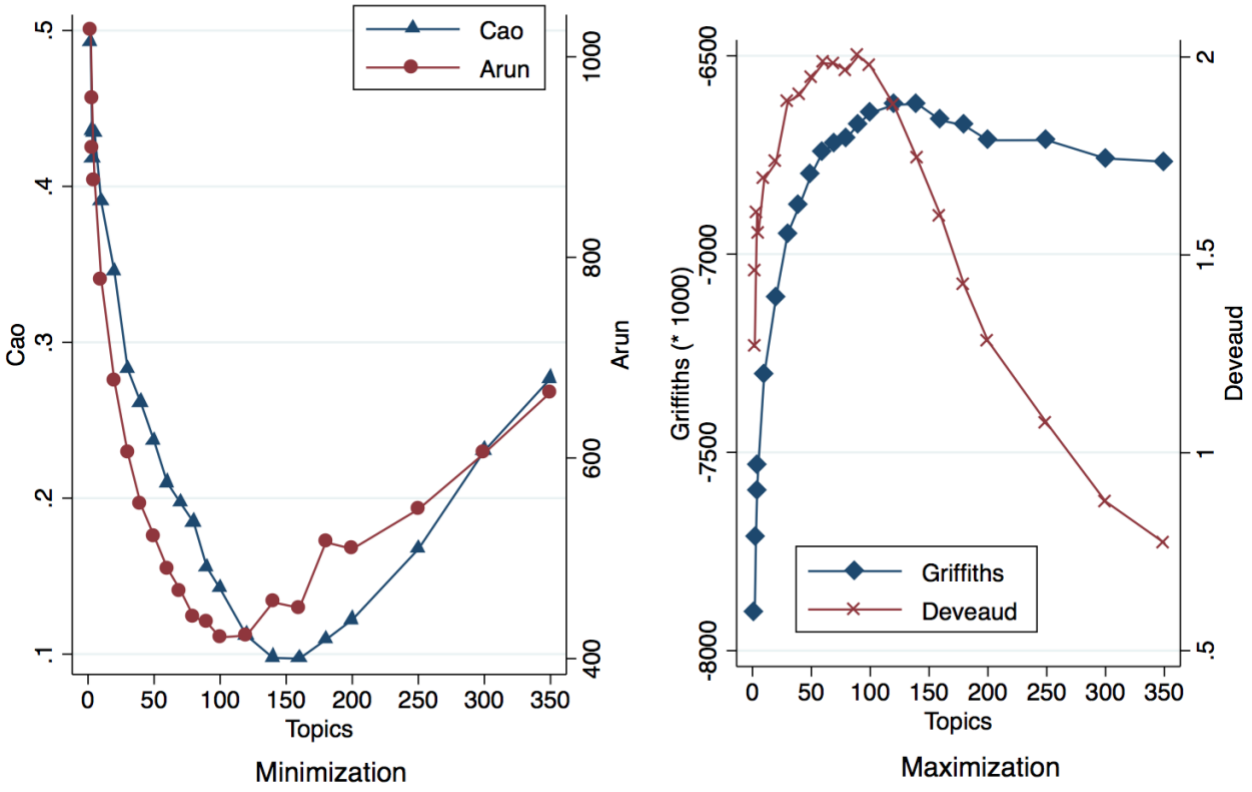

Fortunately, all methods have statistical profiles that indicate that the process of model selection is informative, as the models perform worse after a certain point, indicating that they avoid the problem of overfitting (Griffiths and Steyvers 2004:5231–32). They also converge in agreement that the ideal number of topics for the UNSC resolution dataset is between around 100 and 130 topics, with one method suggesting that the ideal number of topics is closer to 150. Despite the finding that 100 or more topics would be most appropriate, it is not reasonable to conclude that the analysis cannot benefit from fitting models with a smaller number of topics. Sometimes, models that are "better fits" statistically perform relatively poorly from the perspective of the theoretical concerns of the researcher, such that it is standard practice that "several different numbers of topics are fitted and the optimal number is determined in a data-driven way." (Grün

and Hornik 2011) The purpose of this paper is to track changes in abstract concerns over time, and while a model that gives us individual topics of "North Korea," "the Cyprus conflict," and "the former Yugoslavia," may be the "best fit" from a statistical perspective, using such a model would provide limited insight.

In order to judge models, researchers have tended to use two methods: the cohesiveness of topics, or word intrusion, and how closely the topic composition of individual documents matches what would be expected from reading them, or topic intrusion (Chang and Blei 2009:81; Ponweiser 2012:26). This article therefore uses these methods in order to start with the smallest possible number of topics and see whether underlying concepts of interest can be found. Because of the focus of this article, a desirable model will be coherent and create topics on the basis of abstract qualities and not region or any other more concrete metric. Once that is done, however, I also create a model with $k = 110$. This is near the lower range of the estimates given by the four methods used to select the correct number of topics in Figure 2, chosen on the general principle of parsimony, or the idea that simpler models are preferred, all else being equal (Hansen and Yu 2001:96; Posanda and Buckley 2004). The $k = 110$ model serves as a robustness check to investigate whether similar topics of interest emerge in models that are better performing statistically and also to gain more refined insight into the change in data over time.

The content of each topic is determined by finding the terms with the highest β values for each topic within a model. The symbol β represents, for each word, the base 2 log of the probability of that word being in that particular topic, divided by the average

probability of that word appearing across the other topics. In formal terms, this means that

$$\beta_i^t = log_2(\frac{p(i|t)(k-1)}{\sum_{j=1}^{k-1} p(i|j)}) \quad (1)$$

Where β_i^t is how strongly word *i* is associated with topic *t*, all other topics in the model are represented by *j*, and *k* is equal to the total number of topics in the model. Only words that appear with a probability > .001 in any of the topics in the model are considered. All models put the entire corpus through the Gibbs sampling process 6,000 times after the initial random allocation of terms to topic. This entire process is repeated 4 more times for each model with a different *k* value, and only the model with the highest overall posterior likelihood is maintained (Grün and Hornik 2011).

The models are also used to find the θ distribution for each document. If we allow θ_d^t to represent the percentage of document *d* devoted to topic *t*, we end up with the following equation.

$$\theta_{\mathrm{d}}^{\mathrm{t}} = \frac{(\text{number of words in document d belonging to topic t})}{(\text{total number of words in document d})} \quad (2)$$

The values β and θ can be used to judge word intrusion and topic intrusion, as the terms with the highest β values within a topic can indicate whether a topic has coherence, at which point that topic can be given a name. Once that is done, θ distributions for specific documents can be used in order to determine whether the topics most closely associated with any particular document actually appear to reflect its content. Finally, θ

distributions for individual documents over time can be used to measure how the focus of the UNSC has changed over the course of its history.

Results

As the number of topics increases, the model classifies terms less by abstract qualities and more according to region.

Metatopic	Punitive	War	Humanitarian
2	Punitive		Humanitarian
3	Punitive	War	Humanitarian
4	*WMD/ Terror* *Other*	War	Humanitarian
5	*States* *Individual*	*Central/West Africa* *Other*	Humanitarian
6	*WMD/ Diamond* *Police*	*Central/West Africa* *Europe* *Other*	Humanitarian
7	*WMD/ Diamond* *Police*	*Interstate* *Intrastate* *NATO*	*West Africa* *Other*
8	*WMD* *Police* *Piracy/ Terror*	*South Africa/Israel* *West/ Central Africa* *Middle East* *Other*	Humanitarian

Table 1. Topics resulting from models k = 2 through k = 8. Results show that, at the most abstract level, UNSC Resolutions can be classified according to the concepts of *Punitive*, *War*, and *Humanitarian*. Models with more topics simply divide these metatopics into finer topics, often based on region.

Table 1 shows the topics that appear depending on the value of *k*, representing number of topics, as judged by the terms with the highest β values for each topic. I call metatopics those topics that are broken down into more specific components as *k* increases. We see that when $k = 3$, the first topic is *Punitive*, which involves the international community sanctioning a bad actor, including terrorist organizations and

states undertaking WMD programs. *War* refers to an armed contest between two or more sides, whether states or factions within a country. Finally, the topic *Humanitarian* is associated with terms focused on making life better for individuals within a country, particularly in a peacekeeping environment. Running the same models but with more topics leads to the emergence of subtopics that are based on region or more specific versions of the main metatopics. For example, when $k = 4$, *Punitive* is broken down into one topic addressing WMDs and terrorism and another related to other issues. Consistent with Figure 2, Table 1 shows the emergence of region as a categorization variable after $k = 4$. Because this inquiry is not interested in the regional focuses of the UNSC, I start by using the $k = 3$ model, which contains the abstract metatopics that are still present in narrower forms as k increases. Figure 3 shows the 40 terms with the highest β values in the three topics created by the model. See the Online Appendix for the terms with the highest β values for each topic shown in Table 1.

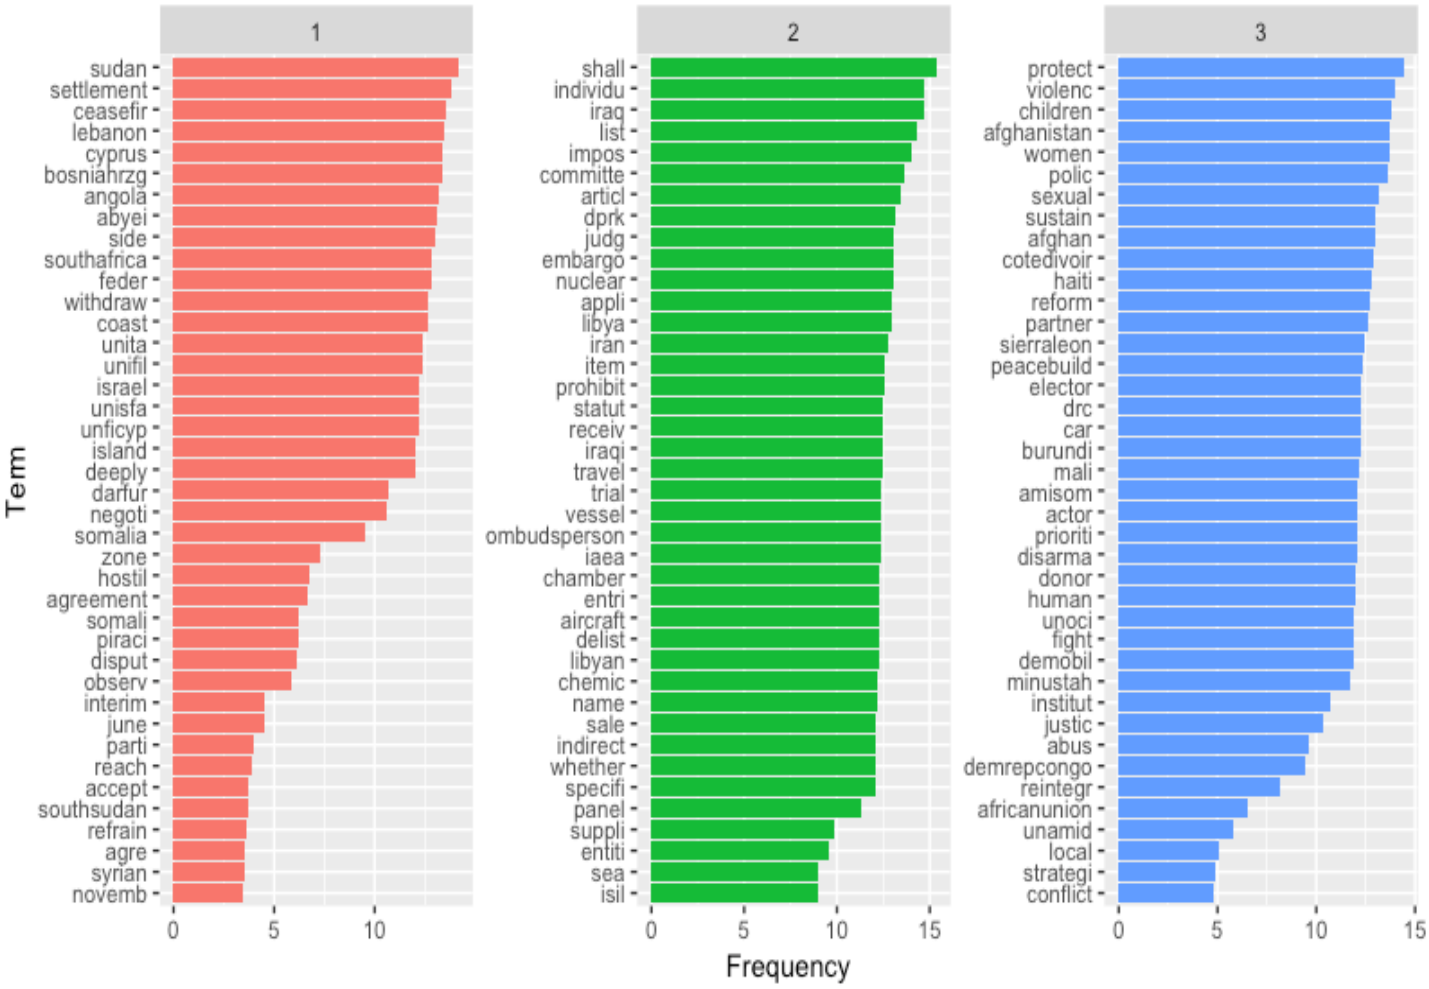

Figure 3. β values for 40 terms most closely associated with each topic.

The patterns are clear. The first topic, *War*, has as its most distinctive terms those that relate to armed conflict. This includes words referring to actions that parties are to take (e.g., *ceasefir*, *withdraw*, *refrain*), locations that may be in dispute (*island*, *zone*) or the combatants themselves (*parti*, *side*). The topic also encompasses processes undertaken to deal with the aftermath of conflicts (*agreement*, *judg*). The countries in that category are ones that have been involved in conflict, whether civil wars (*bosniahrzg*, *somalia*) or international disputes (*israel*, *southafrica*, *somalia*). Topic 3 also contains

terms that refer to war zones, but the abstract words that are most characteristic of that topic generally refer to humanitarian issues that have little direct bearing on the battlefield or international security (e.g., *protect, children, women, sexual*). We can call this topic *Humanitarian*, and its rise over time can be taken to represent the Humanitarian Turn at the Security Council. It also includes many terms relating to peacekeeping and rebuilding and reforming societies in the aftermath of conflict (*polic, disarma, demobil, reintegr*). One may be tempted to see Topic 3 as relating to peacekeeping rather than humanitarian issues, yet the most frequent terms in Topic 1, *War*, include references to several countries that have had UN peacekeeping missions such as Somalia and Cyprus and the missions themselves such as UNFICYP and UNISFA. Thus, Topics 1 and 3 are not distinguished by the abstract categories of peacekeeping and non-peacekeeping. Rather, Topic 1 deals directly with conflict and negotiations between parties, while Topic 3 is mostly characterized by issues relating to improving the lives of people within a state.

Finally, Topic 2, or *Punitive*, is composed partly of a rogues' gallery of outcasts of international society in the postwar era, whether they have been state (*dprk, iran, iraq, libya*) or non-state (*isil*) actors. It includes the policy areas that the international community takes issue with (*chemic, nuclear*), along with the actions that the UNSC takes to punish or deter these bad actors (*embargo, prohibit*). The two most distinctive terms within the topic are *shall*, referring to concrete actions that the Security Council or others are required or resolved to take, and *individu*, reflecting the growing tendency of the institution to penalize individual actors. If we expand our focus beyond Figure 2, notable terms among those with the 100 highest β values for each topic are *area, border,*

delay, *extend*, and *sovereignty* for *War*; *alqaida*, *asset*, *counterterror*, *materiel*, and *transfer* for *Punitive*; and *civilian*, *crime*, *develop*, *transpar*, and *humanitarian* for *Humanitarian*.

The categories created thus pass the test of topic coherence, meaning that they group words together in ways that make sense to an observer (Foltz, Kintsch, and Landauer 1998). Another way to investigate whether the model returns sensible results is to plot individual resolutions along the three dimensions and see whether they fall about where we would expect. This is done in the top triangle in Figure 4, where each point represents a resolution, each of which is placed into one of eight categories indicated by color and shape. Together, the points on the figure represent 460 resolutions, or about an eighth of the total corpus.

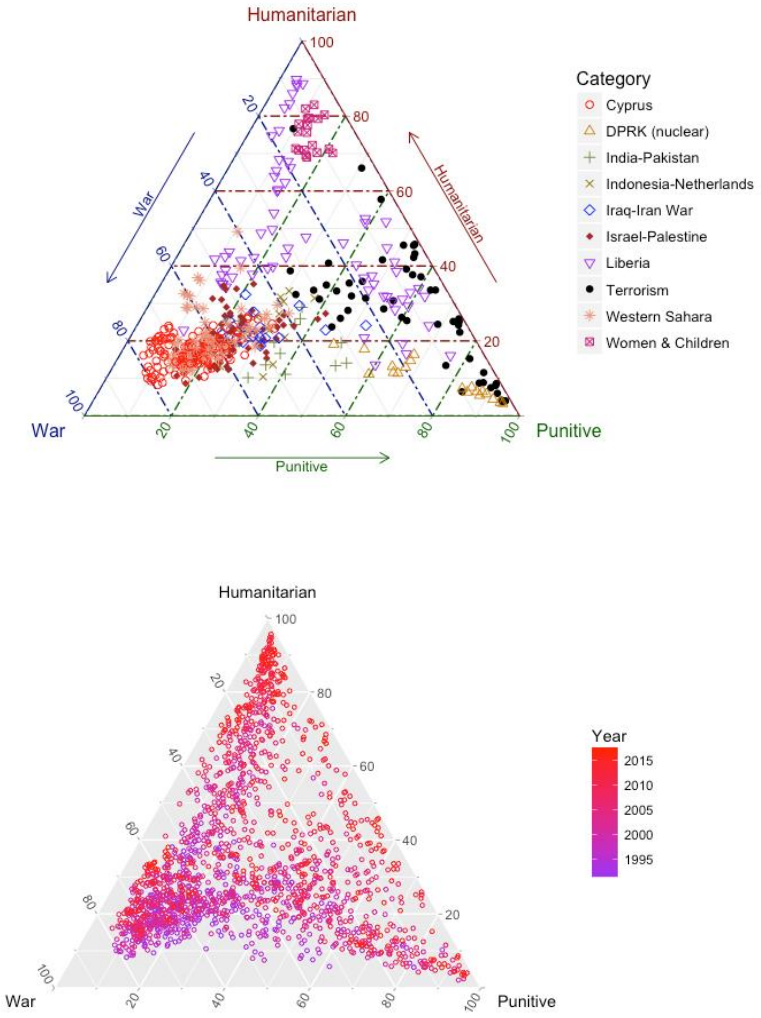

Figure 4. Placement on 3D scale for selected categories of resolutions (top graphic). Placement on 3D scale for all resolutions, 1992-2017 (bottom).

In the top fgraphic, the bottom left corner includes the various conflicts, the bottom right features resolutions related to terrorism and North Korea, and those focusing on women and children are near the top. UNSC resolutions focused on Liberia include components of all three topics, as the UN has engaged in peacekeeping in that country in the midst of a civil war and also sanctioned human right violators. Finally, the Indonesian struggle for independence lies between *War* and *Punitive*, as the UNSC occasionally

pressured the Netherlands in language that is usually reserved for more punitive documents, such as Resolution 67, wherein the Council demanded, among other requests, that the Netherlands release political prisoners and transfer authority over to the Indonesians. The bottom triangle in Figure 4 shows all resolutions from 1992-2017 and reveals changes in topic frequency over time.

The three-topic model thus works to provide coherent topics and classify individual resolutions in ways that make sense. We can therefore move on to investigate how the frequency of each topic per resolution has changed over time in Figure 5, which reveals in various models the average portion of words per document in each year derived from the underlying topics.

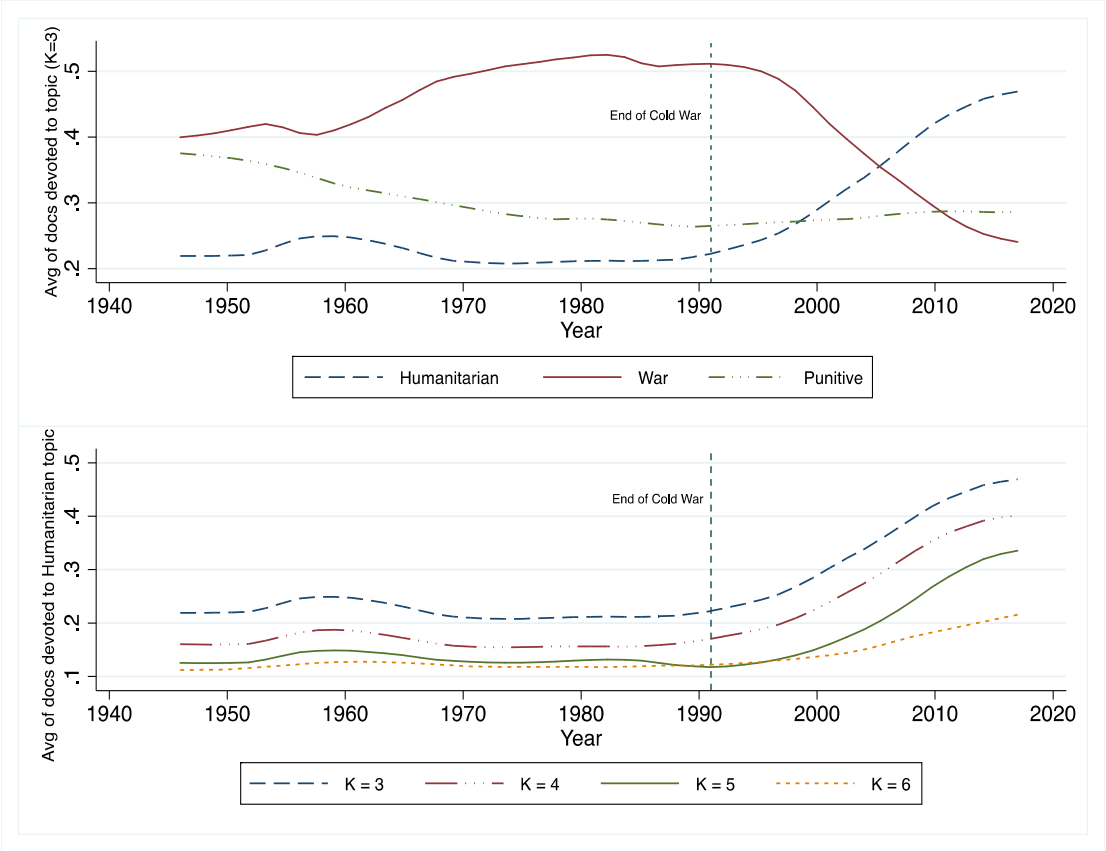

Figure 5. The top graph shows the k=3 model, for each year the average number of words per resolution addressing each topic. The bottom graph shows how the prevalence of the topic *Humanitarian* changes as *k* is increased.

In the three-topic model, *Punitive* stays about constant, dropping off a bit in the first few decades but then resurfacing with the emergence of terrorism and the North Korean nuclear problem as issues. From the founding of the UN until the first decade of the twenty-first century, *War* is the most prominent topic at the Security Council. Finally, *Humanitarian* starts low and remains constant, but then shoots up immediately after the fall of the Berlin Wall. The bottom graph of Figure 5 is presented as a robustness check, showing how the topic *Humanitarian* varies over time when *k* is set to different values. Regardless of the model specification, we see the exact same pattern, with *Humanitarian* taking off around the beginning of the 1990s.

Arguably, the results observed might be explained by changes in power dynamics that made cooperation between superpowers easier. Yet, the nature of the changes in topic distributions over time appears to contradict this simple theory. The topics *War* and *Punitive* address issues that are likely to be of greatest concern in a superpower rivalry: which state and non-state actors to sanction and how to settle disputes between states. Great powers often disagree about whether a bad actor should be punished if one of the powers is allied with the norm violator, as can be seen in efforts to hold the Syrian government accountable for atrocities committed since the beginning of the civil war in that country. They also may take different sides in a conflict, seeing any potential action in zero-sum terms. Yet when the Cold War ended, the most contentious issues began to receive less attention relative to what seem to be less controversial topics. While it is true that the collapse of the Berlin Wall coincided with an increase in peacekeeping, the fundamental nature of peacekeeping missions changed. In the few cases that ended before 1989, "the primary purpose was less to prevent the resumption of war than to contain the conflict to prevent direct superpower interventions." (Fortna 2008:4) Although UN peacekeepers were sent to Cyprus and the Republic of Congo in the 1960s, the resolutions authorizing and supporting those missions are closer in content to the resolutions surrounding interstate disputes than they are to modern peacekeeping resolutions.

Furthermore, if we are tempted to explain the increase in the topic *Humanitarian* by pointing to better relations between the US and Russia, then we may have reason to expect the trends that began at the end of the Cold War to have reversed themselves by now due to deteriorating relations between those powers. The two countries have moved

further apart since the mid-1990s, as can be seen in differences regarding NATO expansion and conflicts in the former Yugoslavia and the Ukraine, and is reflected in voting patterns at the General Assembly (Bailey, Strezhnev, and Voeten 2017). Yet this increase in superpower tension does not seem to have mitigated or reversed trends related to increasing Security Council output, particularly regarding the *Humanitarian* topic, that were set in motion with the collapse of the Soviet Union. This supports the view that international institutions are often characterized by a great deal of inertia that can only be disrupted by sudden changes.

It is important to note that after the Cold War, it cannot be said that the Security Council began to pay less attention to *War* and *Punitive* in absolute terms. In fact, if measured by total number of words or resolution devoted to each topic, activity has universally increased over the last few decades. Figure 6 shows the total number of resolutions passed and words per document for the period between 1946 and 2017.

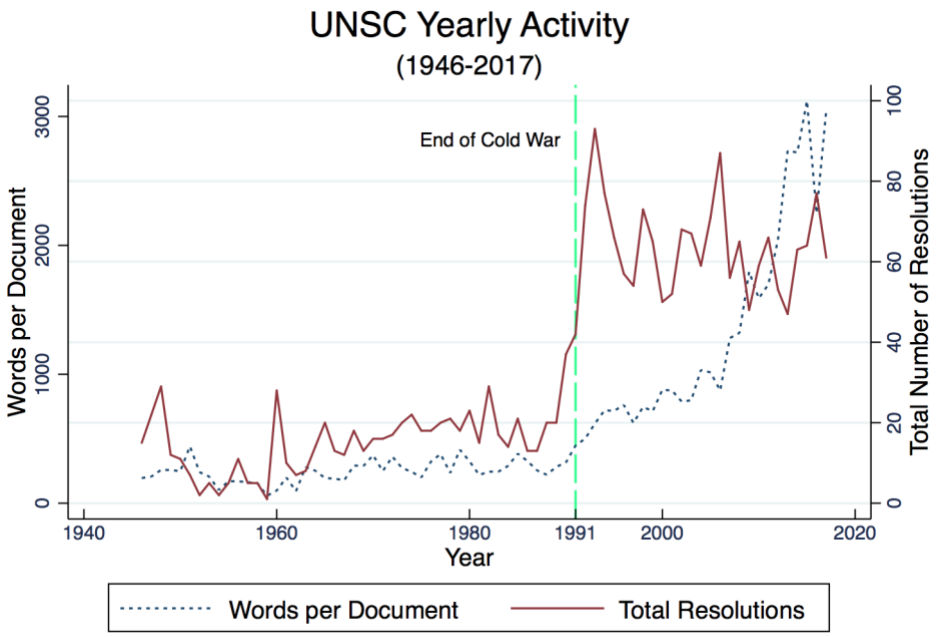

Figure 6

Once again, the end of the Cold War stands out in importance. After that point, the Security Council started passing a larger number of resolutions with a higher average word count. Total activity was stable from the founding of the United Nations up to the end of the Cold War. Interestingly, the number of resolutions has stayed relatively constant since its initial uptick, while words per document has continued to rise to the present day. This disparity must remain unexplained, although it may reflect a more general law about bureaucracies increasing their output over time (Wilson 1989:339–49). The increase in Security Council activity makes the Humanitarian Turn all the more impressive; while there was an increase in relative attention paid to humanitarian issues that can be directly traced to around 1991, in absolute terms the spike in attention devoted to such concerns was even larger.

Returning to Figure 2, we see that the ideal number of topics in the corpus at the lower range is between 100 and 120. This part selects the mid-point of this range, choosing a model where $k = 110$. As expected, doing so produces a large number of

topics pertaining to specific regions and conflicts. At the same time, we see the appearance of a number of distinct abstract categories, which are as expected more specific than those in the models with smaller numbers of topics. Figure 7 shows the change in topic frequency over time for selected topics in the *k = 110* model. At the top of each graph is the top ten terms that are most associated with the given topic.

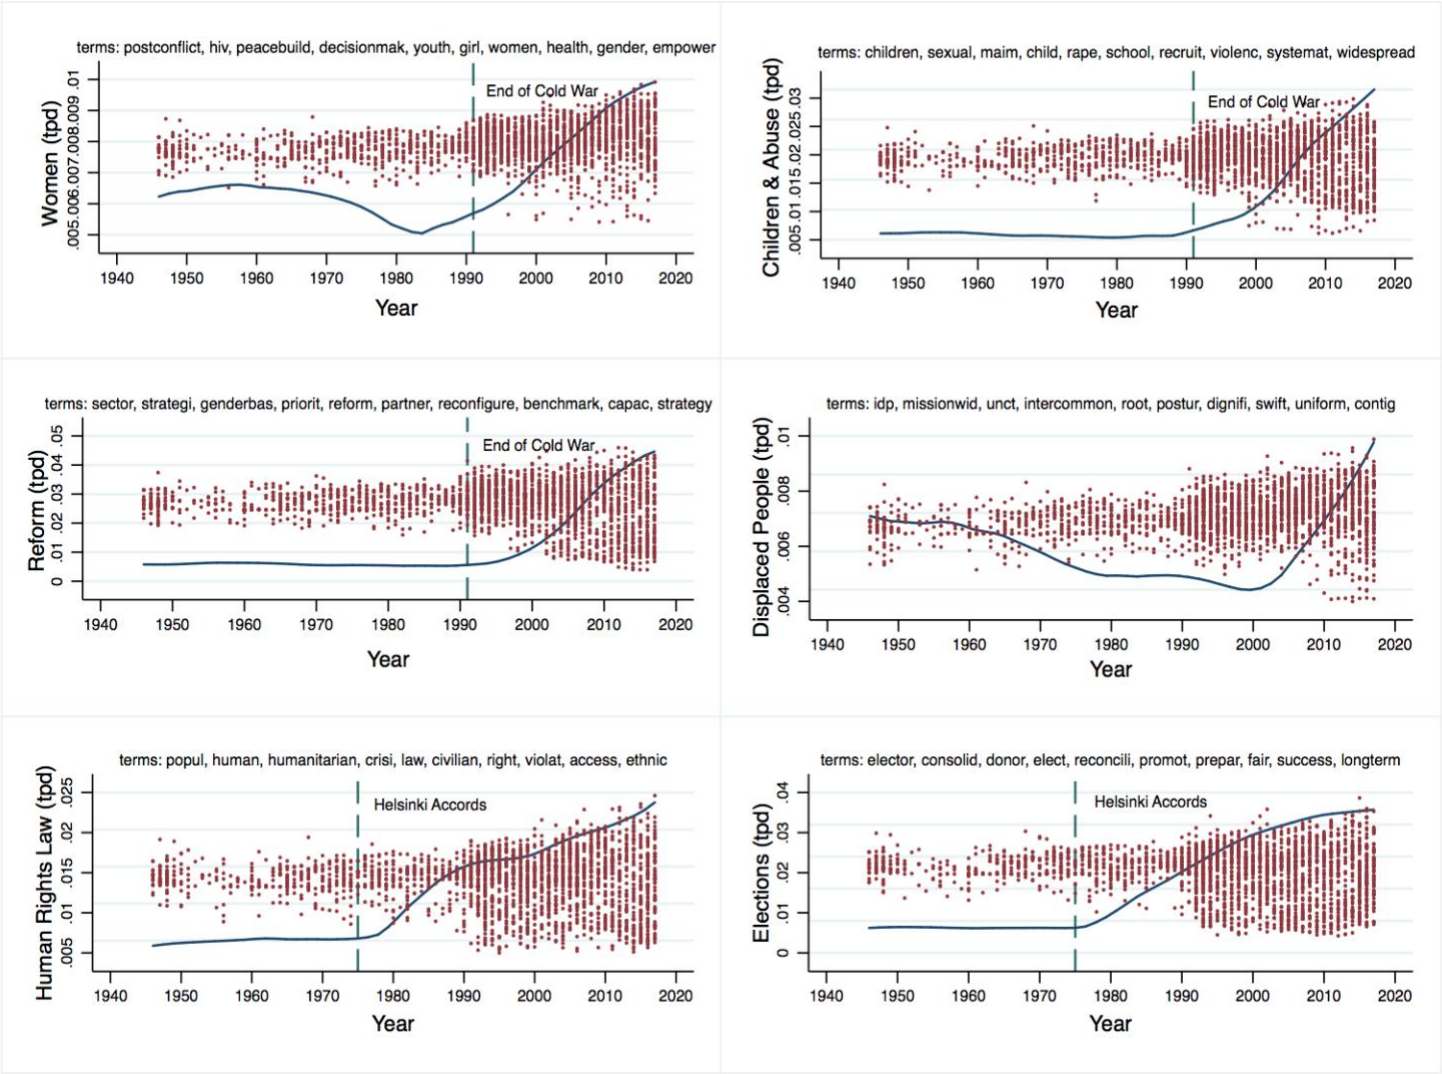

Figure 7. Change in topic focus over time for selected topics, k = 110. Locally smoothed lines are based on the topic per document (tpd) average for the year across all resolutions. Reference points mark the end of the Cold War (1991) or the signing of the Helsinki Accords (1975). Scatter points are based on the negative of the log value of the topic frequency (scale not shown).

We see the clear emergence of categories for *Women, Children & Abuse*, and *Reform*, the last of which contains words typically used in order to set and meet goals within an institutional framework. Again, the end of the Cold War appears important. We also see sharp increases in *Human Rights Law* and *Elections* that correspond to the signing of the Helsinki Accords. Thus, while the Helsinki Accords do not appear to have had an impact from the perspective of the models with lower *k* values, a more fine-grained analysis reveals that they may have been important for increasing attention paid to these two topics. This is consistent with assessments about the central role that the Helsinki Accords played in motivating human rights reforms across the world.

Elections in general refers to UN action undertaken and aspirations put forth in peacekeeping mission; resolutions before the 1970s focused relatively little on the form of governments that states hosting missions should have. Because the changes that we see come near the end of the 1970s, rather than the start of the decade, we have evidence that the Helsinki Accords themselves were the cause of greater focus on human rights law and elections, rather than both being driven by the same underlying historical process. The change that resulted from the Helsinki Accords may have set the stage for the more comprehensive Humanitarian Turn that would arrive a decade and a half later.

As in the models with smaller numbers of topics, it is important to emphasize that Figure 7 addresses topic per document (TPD) frequency. Thus, considering the already noted increase in Security Council activity, the rise of the humanitarian topics is even more striking. Taking into account increased UNSC activity also reveals new topics that took off after the signing of the Helsinki Accords. Figure 8 shows trends regarding the topics *Human Trafficking* and *Media & Medicine*, the latter of which refers to groups that

are expected to be protected within conflict zones. The top graphs shows TPD, or relative frequency, while the bottom graphs takes into account the increasing number of resolutions over time. No account is taken of changes in words per document, as doing so is not necessary to see the trends. The bottom graphs in Figure 8 thus show the total number of resolutions devoted to each topic per year.

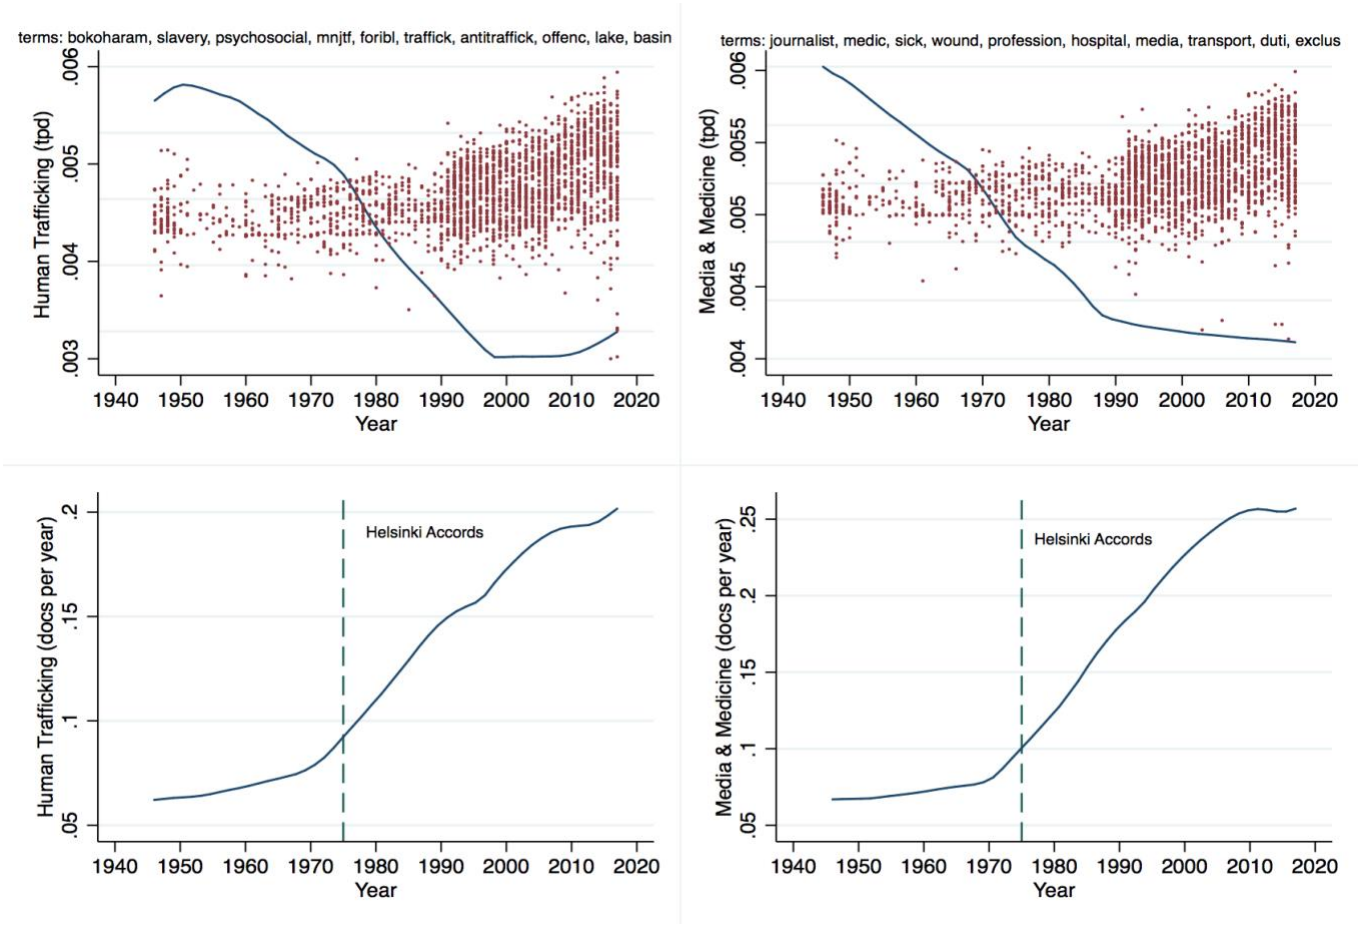

Figure 8. The graphics in the top row follow the pattern of Figure 7. Although the tpd of the two topics decrease over time, the bottom row corrects for the increase in documents produced over time and on the y-axis shows the number of documents in total devoted to the topic per year, or the TPD multiplied by the total number of resolutions per year.

Because TPD is a zero sum-measure, the increase in the prominence of the topics shown in Figure 7 must come at the expense of other topics. This means that using TPD to investigate change in focus over time with regards to *Media & Medicine* and *Human*

Trafficking leads to a misleading picture of the interest of the UNSC in these matters. When we account for the fact that the number of resolutions has increased over the years, we see more focus on these topics that began approximately around the time of the Helsinki Accords. Because these changes appear to have begun before 1975, they can be seen as part of the same process that led to the signing of that treaty. Thus, while focal point events are important, we see some evidence for a trend that was based on a more gradually forming foundation.

Conclusion

Through the use of automated content analysis, this article shows that we can trace changes in the behavior of the Security Council back to the focal point events of the signing of the Helsinki Accords and the end of the Cold War. This project will hopefully be useful to others investigating norm diffusion in international politics. Constructivists stress the importance of norms, and have made progress in tracing the processes through which socialization occurs (Adler-Nissen 2014). They have shown how norms can develop in unexpected ways after they have emerged (Krook and True 2012). This paper argues for the importance of focal point events, particularly the Helsinki Accords and the end of the Cold War. Scholars should continue to explore other events of historical importance in order to test the theory that major public events can inspire new practices, particularly in international institutions.

Text analytic methods can provide a useful service in this area. Those considering delving into the historical record in order to learn about international politics are often cursed with too much information (Trachtenberg 2006:39–45). In studying the UNSC, for example, scholars have access to, among other sources, official documents from states

and the Council itself, minutes of meetings, news reports, press releases of NGOs, and memoirs of individual actors. Automated text analysis can be used to sift through mountains of data, and shine a light on which eras are most worthy of attention. By showing the importance of focal point events, the findings presented here can hopefully make the project of studying norm diffusion in international politics more manageable by suggesting a temporal focus. Researchers who are interested in tracing the Humanitarian Turn at the Security Council but unsure about where to focus their attention would likely be best served by beginning with an investigation into the time periods surrounding the two focal point events highlighted in this article. The findings presented also imply that although insights can be gained by studying the rise of certain topics at the Security Council in isolation, their emergence can be understood more holistically as part of more general changes in international politics (Cohn 2008; Tryggestad 2009).

Perhaps we can understand the Humanitarian Turn in an even broader context. In recent years, scholars have made strong arguments that humanity has seen moral progress (Hathaway and Shapiro 2017; Pinker 2011, 2019). The support for this view is provided in the form of quantitative evidence on the decline of violence and qualitative work on changing practice over time. Automated text analysis provides yet another method to understand these kinds of historical processes. One might always suspect that violence and other undesirable practices have declined due to economic, technological, or other material reasons (Posner 2009; Waltz 1990). Text, however, is perhaps the most direct indicator of what people are thinking in any given time and place. Using it to make empirical arguments is difficult, however, because, as mentioned above, there is often too much information to meaningfully review, along with inherent difficulties in

interpretation. Here automated text analysis can supplement historical inquiry by processing large amounts of data and finding underlying patterns. When the results match what we suspect from the historical record, such findings can give us more confidence about theories of how international society develops.

Automated text analysis is especially likely to be convincing when it is tied to focal point events that do not have direct consequences in terms of economics or the balance of power (Hathaway and Shapiro 2011). When such shocks are not accompanied by an immediate material influence on international politics, such as was the case with the Helsinki Accords, and are nonetheless followed by changes in practice, we can be confident in their ideational impact. And while the end of the Cold War may have changed the power dynamics of international society, the fact that *Humanitarian* is most associated with concepts that are largely unrelated to great power rivalry indicates that a shift in the logic of appropriateness of relevant actors is the reason why (Müller 2004).

In arguing for the importance of focal point events, this article has implications not only for scholars but also for activists and others who would like to effect change in the practices of states and international organizations. Inertia may be the natural state of politics and bureaucracies, yet once in a while we are faced with events that have the potential to create new social realities. The story behind the Humanitarian Turn at the Security Council suggests that the future is often made be those who are able to seize such opportunities.

References

Adler-Nissen, Rebecca. 2014. "Stigma Management in International Relations: Transgressive Identities, Norms, and Order in International Society." *International Organization* 68(1): 143–76.

Arun, Rajkumar, et al. 2010. "On Finding the Natural Number of Topics with Latent Dirichlet Allocation: Some Observations." In Zaki, M.J., et al. *Advances in Knowledge Discovery and Data Mining: Lecture Notes in Computer Science, vol 6118*. Springer: Berlin: 391–402.

Bailey, Michael A., Anton Strezhnev, and Erik Voeten. 2017. "Estimating Dynamic State Preferences from United Nations Voting Data." *Journal of Conflict Resolution* 61(2): 430–56.

Bass, Gary J. *Freedom's Battle: The Origins of Humanitarian Intervention*. New York: Random House.

Blei, David M., Andrew Y. Ng, and Michael I. Jordan. 2003. "Latent Dirichlet Allocation." *Journal of Machine Learning Research* 3: 993–1022.

Bosco, David L. 2009. *Five to Rule Them All: The UN Security Council and the Making of the Modern World*. New York: Oxford University Press.

Braumoeller, Bear F. 2010. "The Myth of American Isolationism." *Foreign Policy Analysis* 6(4): 349–71.

Cao, Juan, et al. 2009. "A Density-Based Method for Adaptive LDA Model Selection." *Neurocomputing* 72(7): 1775–81.

Chang, Jonathan, and David M. Blei. 2009. “Relational Topic Models for Document Networks.” In *Proceedings of the 12th International Conference on Artificial Intelligence and Statistics*: 81–88.

Checkel, Jeffrey T. 1997. *Ideas and International Political Change: Soviet/Russian Behavior and the End of the Cold War*. New Haven, CT: Yale University Press.

Checkel, Jeffrey T. 2005. “International Institutions and Socialization in Europe: Introduction and Framework." *International Organization* 59(4): 801–26.

Cohn, Carol. 2008. “Mainstreaming Gender in UN Security Policy: A Path to Political Transformation?” In Rai, Shirin, and Georgina Waylen, eds. *Global Governance*. Palgrave Macmillan: London: 185–206.

Cortell, Andrew P., and James W. Davis, Jr. 2000. “Understanding the Domestic Impact of International Norms: A Research Agenda.” *International Studies Review* 2(1): 65–87.

Diermeier, Daniel, et al. 2012. “Language and Ideology in Congress.” *British Journal of Political Science* 42(1): 31–55.

Denhert, Gunter. 2013. “The Polish Opposition, the Crisis of the Gierek Era, and the Helsinki Process.” In Eckel, Jan and Samuel Moyn, eds. *The Breakthrough: Human Rights in the 1970s*. Philadelphia: University of Pennsylvania Press: 166–85.

Deveaud, Romain, Eric SanJuan, and Patrice Bellot. 2014. “Accurate and Effective Latent Concept Modeling for Ad Hoc Information Retrieval.” *Document Numérique* 17(1): 61–84.

Evans, Michael et al. 2007. “Recounting the Courts? Applying Automated Content Analysis to Enhance Empirical Legal Research.” *Journal of Empirical Legal Studies* 4(4): 1007–39.

Finnemore, Martha. 1993. “International Organizations as Teachers of Norms: The United Nations Educational, Scientific, and Cultural Organization and Science Policy.” *International Organization* 47(4): 565–97.

Finnemore, Martha, and Kathryn Sikkink. 1998. “International Norm Dynamics and Political Change." *International Organization* 52(4): 887–917.

Foltz, Peter W., Walter Kintsch, and Thomas K. Landauer. 1998. “The Measurement of Textual Coherence with Latent Semantic Analysis.” *Discourse Processes* 25(2-3): 285–307.

Fortna, Virginia Page. 2008. *Does Peacekeeping Work?: Shaping Belligerents' Choices after Civil War*. Princeton University Press.

Gat, Azar. 2006. *War in Human Civilization*. New York: Oxford University Press.

Gheciu, Alexandra. 2005. “Security Institutions as Agents of Socialization? NATO and the ‘New Europe.’” *International Organization* 59(4): 973–1012.

Goldstein, Joshua S. 2012. *Winning the War on War: The Decline of Armed Conflict Worldwide*. New York: Plume Books.

Grimmer, Justin, and Brandon M. Stewart. 2013. “Text as Data: The Promise and Pitfalls of Automatic Content Analysis Methods for Political Texts.” *Political Analysis* 21(3): 267–97.

Griffiths, Thomas L., and Mark Steyvers. 2004. “Finding Scientific Topics." *Proceedings of the National Academy of Sciences* 101(1): 5228–35.

Grün, Bettina, and Kurt Hornik. 2011. “Topicmodels: An R package for Fitting Topic Models.” *Journal of Statistical Software* 40(13): 1–30.

Hansen, Mark H., and Bin Yu. 2001. “Model Selection and the Principle of Minimum Description Length.” *Journal of the American Statistical Association* 96(454): 746–74.

Hathaway, Oona, and Scott J. Shapiro. 2011. “Outcasting: Enforcement in Domestic and International Law." *Yale Law Journal* 121: 252–349.

Hathaway, Oona, and Scott J. Shapiro. 2017. *The Internationalists: And Their Plan to Outlaw War*. New York: Penguin.

Hopf, Ted. 2010. “The Logic of Habit in International Relations.” *European Journal of International Relations* 16(4): 539–61.

Hurd, Ian. 2008. *After Anarchy: Legitimacy and Power in the United Nations Security Council*. Princeton, NJ: Princeton University Press.

Ikenberry, G. John. 2009. *After Victory: Institutions, Strategic Restraint, and the Rebuilding of Order after Major Wars*. Princeton University Press.

Ikenberry, G. John, and Charles A. Kupchan. 1990. “Socialization and Hegemonic Power." *International Organization* 44(3): 283–315.

Johnston, Alastair Ian. 2001. “Treating International Institutions as Social Environments.” *International Studies Quarterly* 45(4): 487–515.

Johnstone, Ian. 2011. *The Power of Deliberation: International Law, Politics and Organizations*. New York: Oxford University Press.

Joachim, Jutta. 2003. “Framing Issues and Seizing Opportunities: The UN, NGOs, and Women's Rights.” *International Studies Quarterly* 47(2): 247–74.

Kaufmann, Chaim D., and Robert A. Pape. 1999. "Explaining Costly International Moral Action: Britain's Sixty-Year Campaign Against the Atlantic Slave Trade." *International Organization* 53(4): 631–68.

Kim, Hun Joon, and Jason C. Sharman. 2014. "Accounts and Accountability: Corruption, Human Rights, and Individual Accountability Norms." *International Organization* 68(2): 417–48.

Krook, Mona Lena, and Jacqui True. 2012. "Rethinking the Life Cycles of International Norms: The United Nations and the Global Promotion of Gender Equality." *European Journal of International Relations* 18(1): 103–27.

Law, David S. 2016. "Constitutional Archetypes." *Texas Law Review* 95: 153–244.

Legro, Jeffrey W. 1997. "Which Norms Matter? Revisiting the 'Failure' of internationalism." *International Organization* 51(1): 31–63.

Lewis, Jeffrey. 2005. "The Janus Face of Brussels: Socialization and Everyday Decision Making in the European Union." *International Organization* 59(4): 937–71.

Liebermann, Oren. December 18, 2017. "US Stands Alone, Defiant at UN Security Council over Jerusalem." *CNN*. Available at https://www.cnn.com/2017/12/18/middleeast/us-un-security-council-jerusalem/index.html.

March, James G., and Johan P. Olsen. 2011. "The Logic of Appropriateness." In Goodin, Robert E., ed. *The Oxford Handbook of Political Science*. New York: Oxford University Press: 478–97.

Mertus, Julie A. 2010. *The United Nations and Human Rights: A Guide for a New Era*. New York: Routledge.

Moyn, Samuel. 2010. *The Last Utopia: Human Rights in History*. London: Belknap.

Mueller, John. 2009. “War Has Almost Ceased to Exist: An Assessment." *Political Science Quarterly* 124(2): 297–321

Müller, Harald. 2004. “Arguing, Bargaining and All That: Communicative Action, Rationalist Theory and the Logic of Appropriateness in International Relations." *European Journal of International Relations* 10(3): 395–435.

O’Neill, Barry. 2005. *Honor, Symbols, and War*. Ann Arbor, MI: University of Michigan Press.

Park, Susan. 2005. “Norm Diffusion within International Organizations: A Case Study of the World Bank.” *Journal of International Relations and Development* 8(2): 111–41.

Pinker, Steven. 2011. *The Better Angels of Our Nature: Why Violence Has Declined*. New York: Viking.

Pinker, Steven. 2019. *Enlightenment Now: The Case for Reason, Science, Humanism, and Progress*. New York: Penguin.

Ponweiser, Martin. 2012. “Latent Dirichlet Distribution in R.” Thesis, Veinna University of Economics and Business. Available at http://epub.wu.ac.at/3558/1/main.pdf.

Posner, Eric A. 2009. *The Perils of Global Legalism*. University of Chicago Press.

Reinisch, August. 2001. “Developing Human Rights and Humanitarian Law Accountability of the Security Council for the Imposition of Economic Sanctions.” *American Journal of International Law* 95(4): 851–72.

Risse, Thomas, and Kathryn Sikkink. 1999. “The Socialization of International Human Rights Norms into Domestic Practices: Introduction.” In Risse, Thomas, Stephen

C. Ropp, and Kathryn Sikkink, eds. *The Power of Human Rights: International Norms and Domestic Change*. New York: Cambridge University Press: 1–38.

Schweigman, David. 2001. *The Authority of The Security Council Under Chapter VII of The UN Charter: Legal Limits and The Role of the International Court of Justice*. Norwell, MA: Kluwer Law International.

Shanks, Cheryl, Harold K. Jacobson, and Jeffrey H. Kaplan. 1996. "Inertia and Change in the Constellation of International Governmental Organizations, 1981–1992." *International Organization* 50(4): 593–627.

Shepherd, Laura J. 2008. "Power and Authority in the Production of United Nations Security Council Resolution 1325." *International Studies Quarterly* 52(2): 383–404.

Simler, Kevin and Robin Hanson. 2017. *The Elephant in The Brain: Hidden Motives in Everyday Life*. New York: Oxford University Press.

Slapin, Jonathan B., and Sven-Oliver Proksch. 2008. "A Scaling Model for Estimating Time-Series Party Positions from Texts." *American Journal of Political Science* 52(3): 705–22.

Spirling, Arthur. 2012. "US Treaty Making with American Indians: Institutional Change and Relative Power, 1784–1911." *American Journal of Political Science* 56(1): 84–97.

Stewart, Brandon M., and Yuri M. Zhukov. 2009. "Use of Force and Civil–Military Relations in Russia: An Automated Content Analysis." *Small Wars & Insurgencies* 20(2): 319–43.

Teh, Yee W., et al. 2005. “Sharing Clusters among Related groups: Hierarchical Dirichlet Processes.” *Advances in Neural Information Processing Systems*. Available at https://papers.nips.cc/paper/2698-sharing-clusters-among-related-groups-hierarchical-dirichlet-processes.

Teitel, Ruti G. 2001. “Humanity's Law: Rule of Law for the New Global Politics.” *Cornell International Law Journal* 35: 355–88.

Thomas, Daniel C. 2001. *The Helsinki Effect: International Norms, Human Rights, and the Demise of Communism*. Princeton University Press.

Trachtenberg, Marc. 2006. *The Craft of International History: A Guide to Method*. Princeton University Press.

True-Frost, C. Cora. 2007. “The Security Council and Norm Consumption." *NYU Journal of International Law and Policy* 40: 115–217.

Tryggestad, Torunn L. 2009. “Trick or Treat? The UN and Implementation of Security Council Resolution 1325 on Women, Peace, and Security.” *Global Governance* 15(4): 539–57.

Waltz, Kenneth N. 1990. “Nuclear Myths and Political Realities.” *American Political Science Review* 84(3): 730–45.

Widmaier, Wesley W., Mark Blyth, and Leonard Seabrooke. 2007. “Exogenous Shocks or Endogenous Constructions? The Meanings of Wars and Crises.” *International Studies Quarterly* 51(4): 747–59.

Wilson, James Q. 1989. *Bureaucracy: What Government Agencies Do and Why They Do It*. New York: Basic Books.

Xie, Boyi, and Rebecca J. Passonneau. 2012. “Supervised HDP Using Prior Knowledge.” In Bouma, G., et al., eds. *Natural Language Processing and Information Systems: Lecture Notes in Computer Science, vol 7337.* Springer: Berlin: 197–202.

Online Appendix for *Explaining the Development of International Norms: The Humanitarian Turn at the United Nations Security Council*

This Online Appendix has two parts. Part I elaborates on the LDA method used and described in a more streamlined fashion in the original paper. Part II shows the words with the top 10 β values for each of the topics presented in Table 1 of the manuscript.

I. The LDA Process

Within the category of automated content analysis, a main divide is between supervised and unsupervised learning models (Grimmer and Stewart 2013). In a supervised learning model, the researcher begins by inputting classifications of some randomly chosen subset of the data. This might be known through human knowledge and intuition, or through a structured coding process. Then, a model is built, which is used to classify the remainder of the set of documents, referred to as the corpus, or even future documents (Guzella and Caminhas 2009). Researchers do not simply use human coding for the entire corpus because doing so would be too labor intensive. Unsupervised methods are used when the researcher is seeking to find structures in underlying data and does not know what patterns will emerge (Qin, Geng, and Liu 2010). The researcher in effect simply inputs the text, sets certain parameters, and sees what comes out (Spirling 2012).

Building on unigram models and Latent Semantic Analysis, LDA has the advantage of being a mixed-membership topic model, which means that each document is treated as a mixture of topics (Deerwester et al. 1990; Nigam et al. 2000). As an unsupervised machine learning model, it can take in unstructured individual texts as inputs and, as outputs, infer the underlying topics used to create the documents, along

with the word contents of specific documents. In order to begin the process of finding underlying topics in the data, the researcher begins by creating a document-term matrix (DTM), in which each row represents a document and each column represents a term that appears anywhere in the corpus (Haddi, Liu, and Shi 2013). Thus, each entry *i,j* is the number of times term *j* appears in document *i*. Not all terms are included in the column; usually stop words, words that by themselves convey no semantic information such as "the," "but," and "or" are excluded. Words tend also to be stemmed, meaning that related words from the same root are characterized as one term. This would apply, for example to the terms "institution," "institutions," and "institutional."

Gibbs sampling has emerged as the preferred method for conducting LDA (Griffiths and Steyvers 2004). As a hyperparameter, the researcher needs to choose the *k* number of topics. The algorithm then begins by taking each term of each document and randomly assigning it to one of the *k* topics. It also creates a term-topic *n x k* matrix, revealing the count of each word assigned to each topic, and a *d x k* document-topic matrix, which shows the number of words assigned to each topic for each document. Allow w_v to be a word in the vocabulary of the corpus and $w_{u,v}$ to be the particular appearance of that word. For each document the algorithm goes through each word $w_{u,v}$ and assumes that all topic assignments in the entire corpus other than $w_{u,v}$, including other appearances of w_v, are correct. Then, $w_{u,v}$ may be reassigned to another topic. Whether it is reassigned to a particular topic *z* depends on the probability of any word other than $w_{u,v}$ being assigned to topic *z* in document *d* and the probability of w_v appearing in topic *z* across all documents in the corpus, again excluding $w_{u,v}$. In intuitive terms, a particular word is more likely to be assigned to a particular topic if that topic is prevalent in the

document in which it is located, and if the word has a tendency to be assigned to that topic when it appears elsewhere in the corpus. Upon completing this process of for all words in document *d*, the topic probabilities for document *d* are estimated again. The process repeats itself across the entire corpus, and after a set number of iterations, the model stabilizes with each word $w_{u,v}$ assigned to a topic. At that point, it is relatively simple to calculate the topic distribution, θ_d, of each document, and per-topic word distribution of each topic *z*, or z_n, the main outcome of interest for most researchers.

With regards to combining words into terms in the manuscript, it is important to note that combining the names of proper nouns comes first. For example, "Democratic Republic of the Congo" and "Republic of Congo" become individual terms before "Federal Democratic Republic," "Democratic Republic of," and "Republic of" are eliminated. The "Federal Republic of Ethiopia" can then be shortened to "Ethiopia," because it does not need to be differentiated from any other Ethiopia. This prevents the algorithm from classifying parts of names of countries in the same topic together based on coincidental or historically based similarities in their names.

II. Topics Based on *k*

The concept of topic coherence is key to LDA being a useful method (Foltz, Kintsch, and Landauer 1998). The key point for the purposes of this paper is that we should not expect extremely different results based on the number of topics chosen. Table 1 in the main manuscript makes the point that the metatopics of *Humanitarian*, *Punitive*, and *War* are robust regardless of model specifications. This section of the appendix presents the direct evidence for this claim, showing the terms with the top 10 β values for each topic alluded to in the paper.

2-Topic Model

Humanitarian	arm conflict right personnel human civilian protect somalia violenc children
Punitive	committe shall iraq list annex submit tribun present settlement day

3-Topic Models

War	sudan settlement ceasefir lebanon cyprus bosniaandherzegovina angola abyei side southafrica
Punitive	shall individu iraq

	list impos committe articl dprk judg embargo
Humanitarian	protect violenc children afghanistan women polic sexual sustain afghan cotedivoir

4-Topic Models

Punitive (WMD/Terror)	committe individu iraq impos sanction dprk nuclear libya iran item
Punitive (Other)	tribun articl piraci yugoslavia angola judg southafrica sea mr

	coast
Humanitarian	children women afghan cotedivoir haiti partner sierraleon peacebuild drc car
War	sudan southsudan darfur cyprus abyei side unamid hostil unifil monuc

5-Topic Models

Punitive (States)	iraq tribun articl judg nuclear southafrica iran trial chamber southernrhodesia
Punitive (Individuals)	list committe suppli panel isil

	individu item export travel ombudsperson
War (Central/West Africa)	liberia cotedivoir haiti sierraleon burundi mali amisom unoci minustah ecowa
War (Other)	darfur cyprus abyei unamid sudan unifil unisfa unficyp island croatia
Humanitarian	children women piraci peacebuild actor coast guineabissau monusco robberi afghan

6-Topic Models

War (Other)	lebanon

	angola judg southafrica iraqi unita unifil lebanes boundari namibia
Punitive (WMD/Diamond)	dprk nuclear item iaea ombudsperson delist sale notifi petroleum inspect
War (Central/West Africa)	liberia cotedivoir haiti sierraleon elector burundi mali unoci minustah guineabissau
War (Europe)	settlement cyprus unficyp island croatia westernsahara disengag eufor georgia unomig
Humanitarian	sudan

	democraticrepublicofcongo southsudan darfur drc car abyei unamid monuc unisfa
Punitive (Police)	somalia afghanistan afghan piraci terror amisom somali isil sea coast

7-Topic Models

Punitive (WMD/Diamond)	committe dprk nuclear list item iaea ombudsperson delist libyan notifi
War (Interstate)	lebanon angola southafrica unita unifil israel lebanes southernrhodesia

	namibia lusaka
War (Intrastate)	somalia cyprus piraci somali feder coast robberi unficyp westernsahara pirat
War (NATO)	tribun afghan afghanistan bosniaandherzegovina yugoslavia judg mr chamber nato unama
Humanitarian (Other)	sudan southsudan abyei unamid darfur unisfa unmiss mediat chad undof
Punitive (Police)	liberia cotedivoir haiti sierraleon burundi mali unoci minustah

	guineabissau unmil
Humanitarian (West Africa)	democraticrepublicofcongo drc monusco congoles cted greatlak monuc counter child recruit

8-Topic Models

Punitive (WMD)	iraq sudan darfur dprk abyei nuclear iran iraqi unisfa missil
War (South Africa/Israel)	southafrica southernrhodesia namibia aggress armistic heard southafrican zambia apartheid admiss
Punitive (Police)	ombudsperson delist notifi exempt

	judg petition chamber adlitem prosecutor list
War (West/Central Africa)	liberia cotedivoir sierraleon burundi mali unoci guineabissau unmil unita timorlest
War (Other)	haiti cyprus unficyp island croatia multin eufor georgia amisom unomig
War (Middle East)	democraticrepublicofcongo monuc unmiss westernsahara lebanes undof minurso envoy unifil chad
Humanitarian	afghan drc women au

	child genderbas unama children monusco postconflict
Punitive (Piracy/Terror)	somalia piraci somali isil sea coast terrorist robberi libyan counterterror

References (Not Included in Manuscript)

Deerwester Scott, et al. 1990. "Indexing by Latent Semantic Analysis." *Journal of the American Society for Information Science* 41(6): 391–407.

Guzella, Thiago S., and Walmir M. Caminhas. 2009. "A Review of Machine Learning Approaches to Spam Filtering." *Expert Systems with Applications* 36(7): 10206–22.

Haddi, Emma, Xiaohui Liu, and Yong Shi. 2013. "The Role of Text Pre-Processing in Sentiment Analysis." *Procedia Computer Science* 17: 26–32.

Nigam, Kamal, et al. 2000. "Text Classification from Labeled and Unlabeled Documents Using EM." *Machine Learning* 39(2-3): 103–34.

Qin, Tao, Xiubo Geng, and Tie-Yan Liu. 2010. "A New Probabilistic Model for Rank Aggregation." *Advances in Neural Information Processing Systems* 23. Available

at https://papers.nips.cc/paper/3906-a-new-probabilistic-model-for-rank-aggregation.